

Smeet Parinbhai Mehta

Los Angeles, CA | (213) 691-6914 | smeetmehta14@gmail.com | <https://www.linkedin.com/in/smeetmehta/>

EDUCATION

University of Southern California (USC)

Aug 2024 - May 2026

Master of Science in Computer Science - Applied Data Science GPA: 3.77

Relevant Coursework: Fundamentals of Data Management, Data Visualization, Machine Learning for Data Science, Data Mining

SRM Institute of Science & Technology (SRMIST)

Jun 2018 - May 2022

Bachelor of Technology in Computer Science and Engineering GPA: 3.88

SKILLS

Programming, Databases & Systems: Python, Java, C++, SQL (PostgreSQL, MySQL), Linux, Git, Docker, REST APIs

ML / GenAI / NLP: LLMs, Multi-Agent Systems, LangChain, Prompt Engineering, RAG, XGBoost, Scikit-learn, NLTK

Data Engineering & Distributed Computing: PySpark (RDD), Apache Hadoop, Distributed Data Pipelines, Database Design, Query Optimization, MongoDB, Snowflake

Analytics & Professional Skills: Data Analysis, Power BI, Modular System Design, Google Gemini API

EXPERIENCE

Lumina AI Health Institute - *AI Engineer Intern*

Los Angeles, USA | May 2025 - Aug 2025

- Architected and deployed a scalable, PostgreSQL-backed multi-modal clinical AI platform integrating lab results, medical imaging, and genomic data, reducing data retrieval latency by consolidating previously siloed sources into a single intelligence layer.
- Engineered, tested, and deployed 8 LangChain agents via Firebase Studio for medication-dosing analysis, enabling automated clinical recommendations and eliminating hours of manual physician lookup workflows.
- Led full formalization of the AI platform using Google Gemini, establishing structured prompt engineering standards that improved response consistency across 3 distinct medical data modalities

Tiger Analytics - *Senior Software Engineer - Analytics*

Bangalore, India | Jun 2023 - Jul 2024

- Owned end-to-end design, deployment, and maintenance of Insights-Pro, a production multi-agent AI system (Web Search, SQL Generation, Validation agents), achieving 90% accuracy on natural-language business queries and boosting generative AI output efficiency by 80%
- Designed a fault-tolerant agent orchestration layer with state management and error-recovery, enabling reliable inter-agent communication across structured and unstructured data sources for enterprise clients.

Tiger Analytics - *Software Engineer*

Bangalore, India | Jun 2022 - Jun 2023

- Built a menu-based price optimization model for a multinational food giant using derivative-free constrained optimization in Python, delivering a 30% average profit increase by encoding country-specific business constraints.
- Refactored a 100K+ line large-scale production codebase into a modular, region-configurable architecture, cutting technical debt and production incidents by 40% and reducing regional feature rollout time by weeks.

Tiger Analytics - *Software Engineer Intern*

Bangalore, India | Aug 2021 - Jun 2022

- Automated region-specific pricing deployments via a standardized constraint-driven pipeline, eliminating manual adjustment overhead across 10+ markets.

PROJECTS

Yelp Recommendation System | *Python, PySpark (RDD), XGBoost, Scikit-learn*

- Built a hybrid collaborative + content-based recommendation system achieving 0.972 RMSE, outperforming established production benchmarks.
- Processed 1.2GB+ of multi-modal data on a distributed Spark RDD pipeline, engineering 45+ features from user metadata, business attributes, and check-in trends in under 3 minutes end-to-end.

LinguaSense - *Multilingual Sentiment Analysis* | *Python, Hugging Face, LLMs, Sentiment Analysis, Multilingual NLP*

- Designed an explainable multilingual sentiment system combining GenAI-powered translation with LLM-based classification, delivering real-time, language-agnostic predictions with natural-language justifications.

ConvoDB - *Natural Language to Query Interface* | *Python, NLTK, RegEx, SQL, MongoDB*

- Built a zero-LLM NL-to-SQL/MongoDB query engine from scratch using NLTK and RegEx, supporting direct database uploads and achieving accurate query construction across structured and semi-structured schemas.

LEADERSHIP

Viterbi School of Engineering, USC

Los Angeles, USA

Course Producer (DSCI 554- Data Visualization)

Jan 2026 - Present

- Manage grading workflows and course delivery for a graduate-level visualization course.

Viterbi Graduate Orientation Leader

Aug 2025 - Present

- Led Fall '25 and Spring '26 Viterbi Graduate Orientation, guiding 1000+ students from 52+ countries to adapt to campus life at USC